

A full-page background image showing an astronaut in a white spacesuit standing on a dark, rocky surface, looking out over a vast, hazy, orange-yellow landscape under a bright, glowing sky. The astronaut is positioned on the right side of the frame, facing left.

Lunar Lander

Krishan Subudhi

Final Project: University of Washington



Goal

Direct the agent to the landing pad as softly and fuel-efficiently as possible.



Algorithms

Atari

Box2D

Classic control

MuJoCo

Robotics

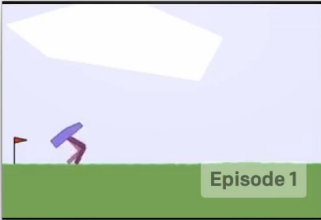
Toy text EASY

Third party environments

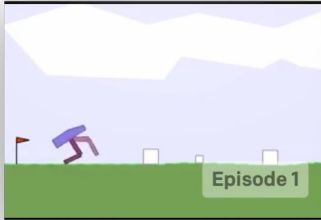
gym.openai.com

Box2D

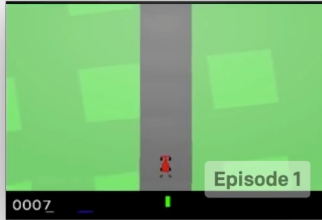
Continuous control tasks in the Box2D simulator.




BipedalWalker-v2
Train a bipedal robot to walk.




BipedalWalkerHardcore-v2
Train a bipedal robot to walk over rough terrain.




CarRacing-v0
Race a car around a track.



LunarLander-v2
Navigate a lander to its landing pad.




LunarLanderContinuous-v2
Navigate a lander to its landing pad.



Setup

- x coordinate
- y coordinate
- horizontal velocity
- Vertical velocity
- angle
- angular velocity
- Is Left leg on ground?
- Is right leg on ground?



1. do nothing
2. fire left orientation engine
3. fire main engine
4. fire right orientation engine

Rewards

- Moving from the top of the screen to the landing pad and coming to rest : +100 to +140
- If the lander crashes : -100
- If it comes to rest : +100
- Each leg with ground contact : +10
- Firing the main engine : -0.3
- Firing the side engine : -0.03



Approach



Why Reinforcement Learning?

We are treating the environment as a blackbox.

We don't know $T(s,a,s')$ and $R(s,a,s')$



Algorithms tried

Random baseline

Q Learning


Approximate Q Learning

Deep Q learning



Random Agent


- Chooses action randomly
- Baseline



Q Learning

- Maintains a dictionary of
 - $\{(state, action)\} : QValue$
- Our states are continuous. Need to discretize.
- Learning
 - $Q_{sample} = R(s, a, s') + \gamma \max_{a'} Q(s', a')$
 - $Q_{t+1} = (1 - \alpha)Q_t + \alpha Q_{sample}$
- Hyper parameters
 - episodes = 10000
 - $\alpha = 0.3$
 - $\gamma = 0.95$
 - $\epsilon =$ step function from 0.5 to 0

LanderQTableAgent
Average rewards



Average steps







Image from [arXiv:2011.11850v1](https://arxiv.org/abs/2011.11850v1) [cs.LG]

Epsilon



Approximate Q Learning

- Q table had a lot of (state, action) pairs.
- Approximate Q Functions
 - $Q(s, a) = \sum_i w_i f_i(s, a) + bias$
- Learning after transition (s, a, r, s')
 - $diff = [R(s, a, s') + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
 - $w_i = w_i + \alpha * diff * f_i(s, a)$
- Hyper parameters
 - episodes = 10000
 - $\alpha = 0.3$
 - $\gamma = 0.95$
 - ϵ = step function from 0.5 to 0
- Findings
 - Converges to lower average reward.
 - Increasing number of exploration steps does not help.
 - Probably underfitting issue.
 - How to verify that algorithm/code is accurate?



Custom game : 1Dtarget

- Goal
 - given a point as target (100) , and a starting state (1) , agent should reach the target as quickly as possible.
- State space
 - agent x position
- Action space
 - Left
 - right
- Reward
 - 1 if agent reaches target

[0 0 0 0 0 0 100 0 0 1 0 0]



Approximate Q Learner on 1DTarget

[0	0	1	0	0	1	0	0	0	0	0	0]
---	---	---	---	---	---	---	---	---	---	---	---	----

[0	0	0	1	0	1	0	0	0	0	0	0]
---	---	---	---	---	---	---	---	---	---	---	---	----

[0	0	0	0	1	1	0	0	0	0	0	0]
---	---	---	---	---	---	---	---	---	---	---	---	----

[0	0	0	0	0	1	0	0	0	0	0	0]
---	---	---	---	---	---	---	---	---	---	---	---	----

Struggle at the edges

[0	0	0	0	0	1	0	0	1	0	0	0]
---	---	---	---	---	---	---	---	---	---	---	---	----

[0	0	0	0	0	1	0	1	0	0	0	0]
---	---	---	---	---	---	---	---	---	---	---	---	----

[0	0	0	0	0	1	0	0	1	0	0	0]
---	---	---	---	---	---	---	---	---	---	---	---	----

Learned Policy:

0: right

1: right

2: right

3: right

4: right

5: right

6: right

7: left

8: left

9: left

10: left



Deep Q Learning



Neural network

Fetches training samples from past memory



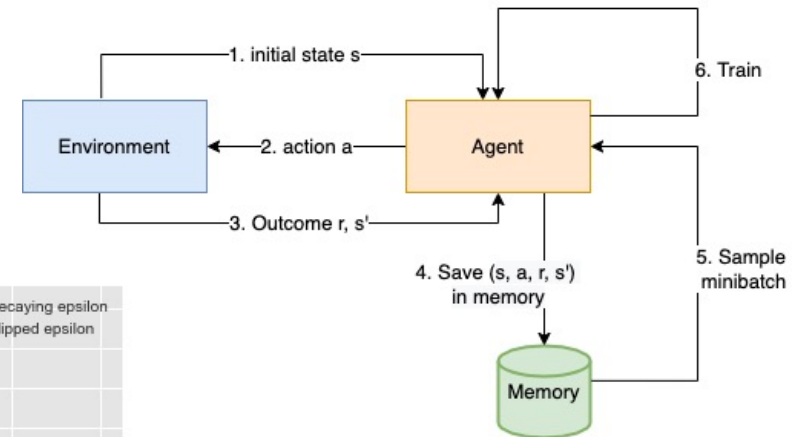
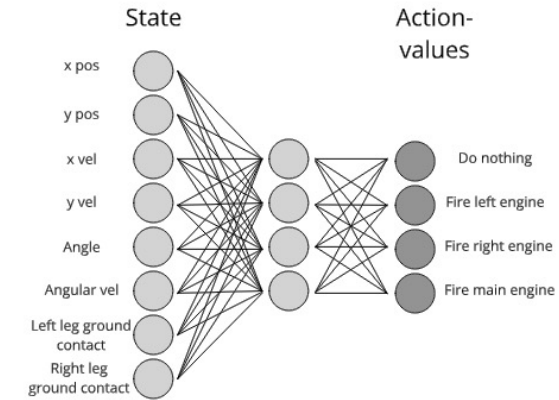
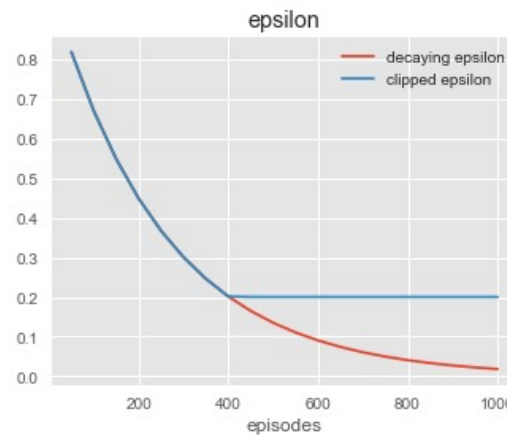

Advantages

Non-linear QValue approximation.
No feature engineering



Challenges

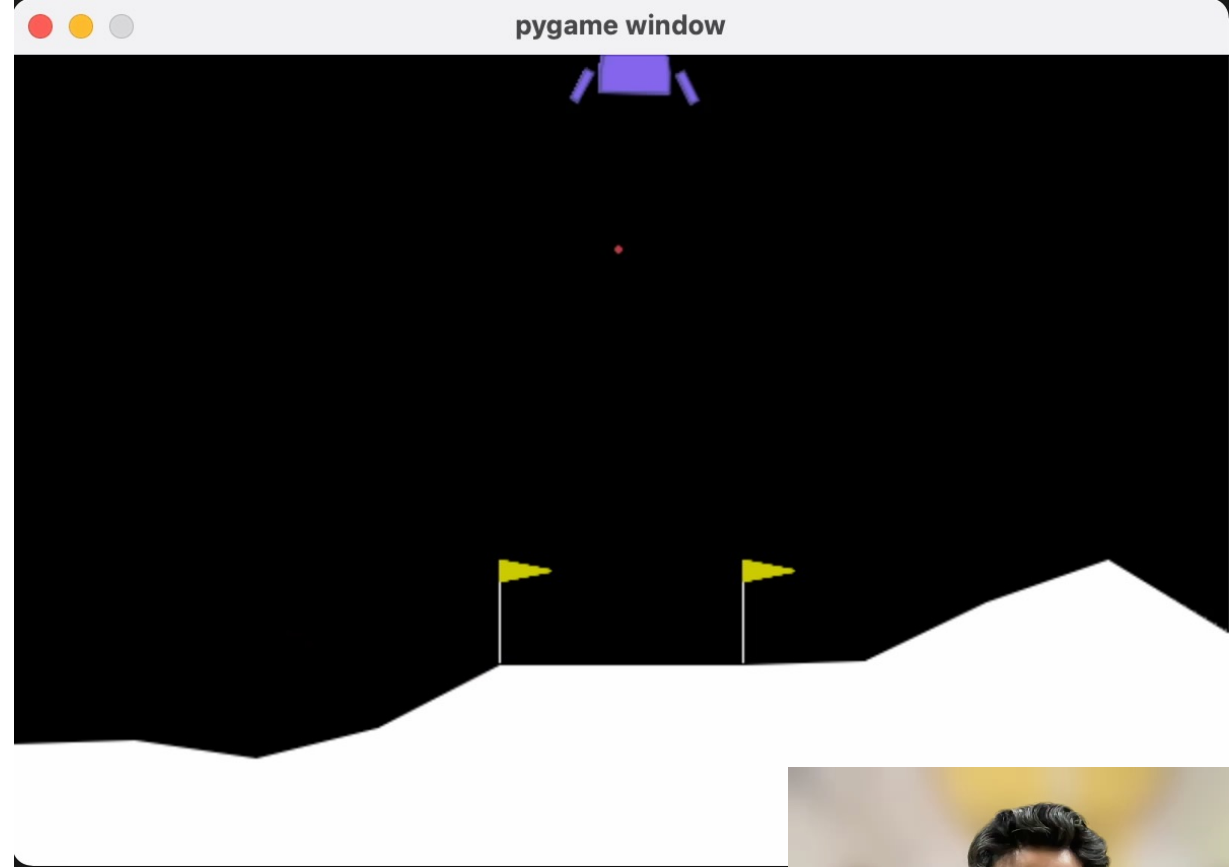

samples are not iid.
Label are derived bellman's equation.



- Maximum posit
- Exploration is in
- Replay quality is
- Low exploration



Before and after training with DQL




steps 328
reward +272.64
action_counts = [14, 83, 110, 121]



Evaluation under different conditions

- `s = env.reset(seed=seed)`
- Seed determines the initial force and terrain.
- Agent was trained with fixed seed.
- Agent performs poorly for random seeds.
- Overfitting?


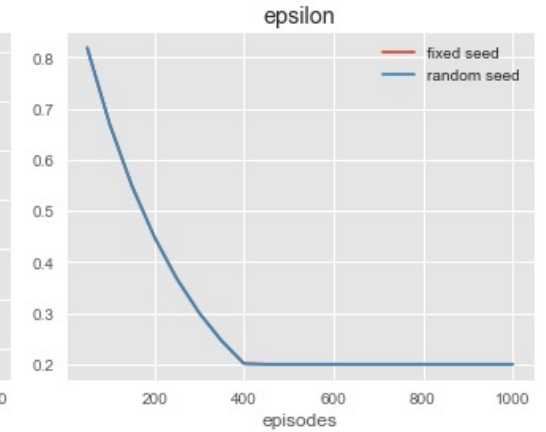
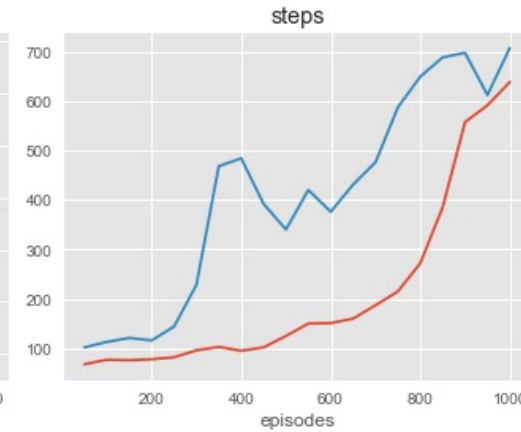


ep 0 steps 108 reward +6.31, action_counts = [0, 28, 5]
ep 1 steps 581 reward +190.21, action_counts = [17, 3]
ep 2 steps 224 reward -57.21, action_counts = [6, 44, 5]
ep 3 steps 215 reward -69.31, action_counts = [7, 39, 5]
ep 4 steps 193 reward +1.95, action_counts = [6, 52, 5]



Training with Random seed

- More diverse states encountered.
- Average reward is low. Needs more hyperparameter tuning.
- ideas?



Learnings

- RL is difficult to debug. Start simple.
 - 1DTarget, Q tables
- Hyper parameter tuning is important.
 - Most important : ϵ
 - Other important hyper params : α , number of episodes, max steps per episode, replay memory size, γ
- Visualization is important
 - Renders, graphs – gather as much info as possible during training
- GPUs don't help in Deep Q Learning. Each step requires inferencing the NN.
- RL is fun!



References

Project source : <https://github.com/krishansubudhi/lunarlander>

Reference paper: [arXiv:2011.11850v1](https://arxiv.org/abs/2011.11850) [cs.LG]

Environment: [OpenAI Gym LunarLanderV2](https://gymnasium.openai.com/environments/lunarlander-v2)

